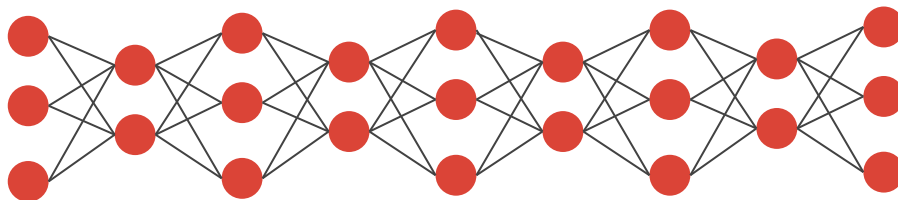


# Deep Learning Fundamentals



Justin Chen

**Boston University** Machine Intelligence Community

# Fluffy Stuff

---

- Motivation
- Application
- Statistical Learning

Intelligent Machines

Patterns in Data

Understand the Mind

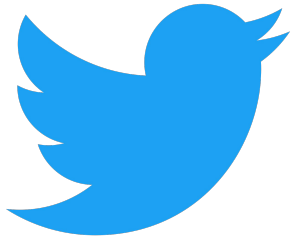
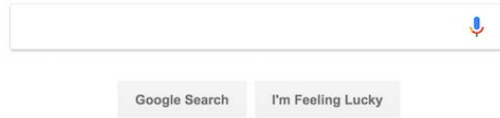
# Application



**NETFLIX**



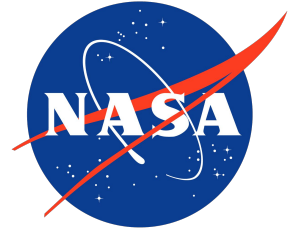
**Google**



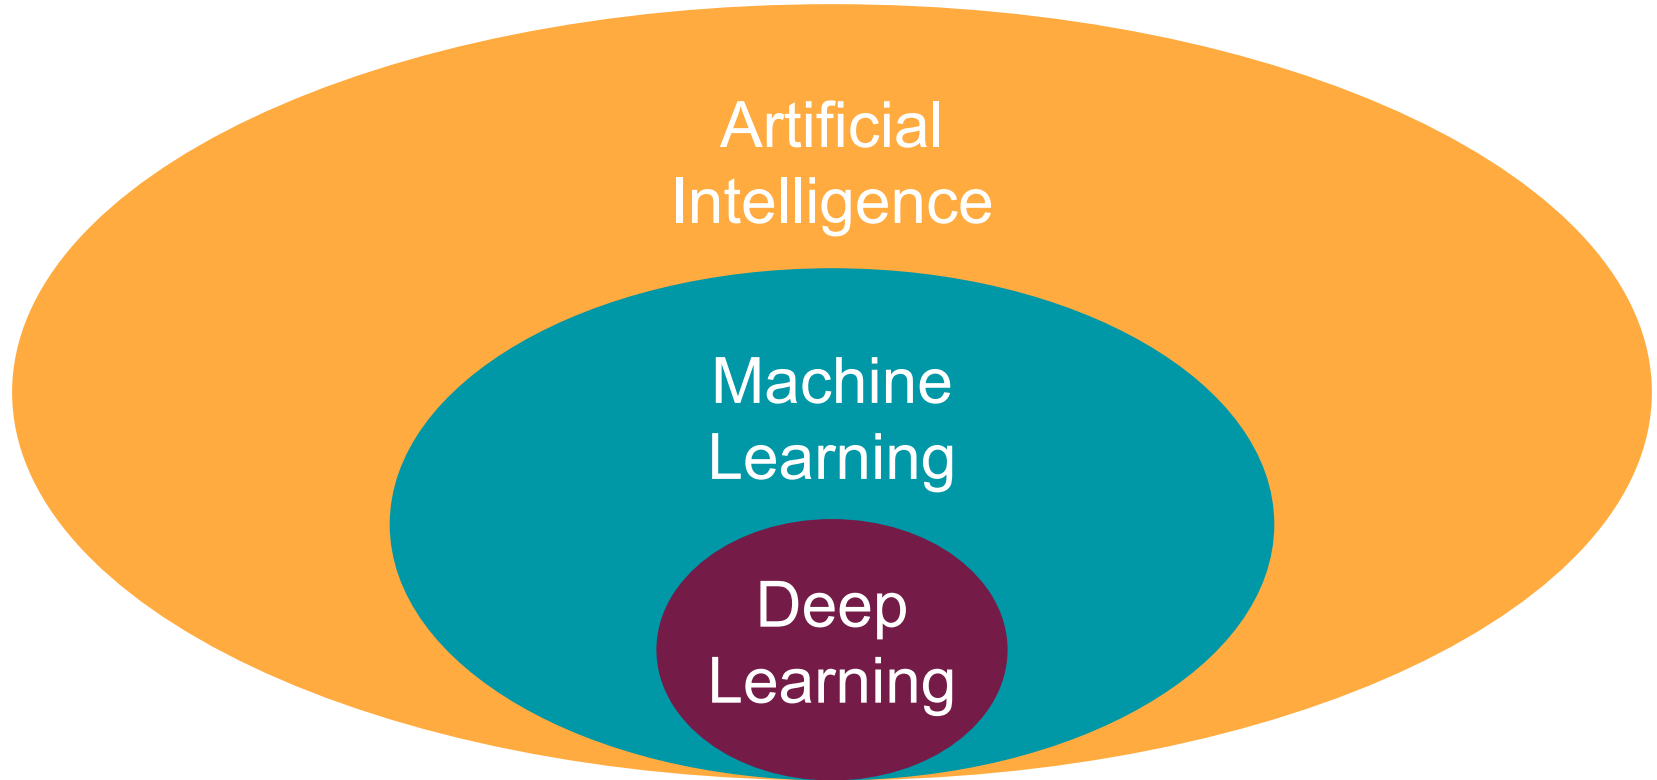
**You Tube**



**amazon**

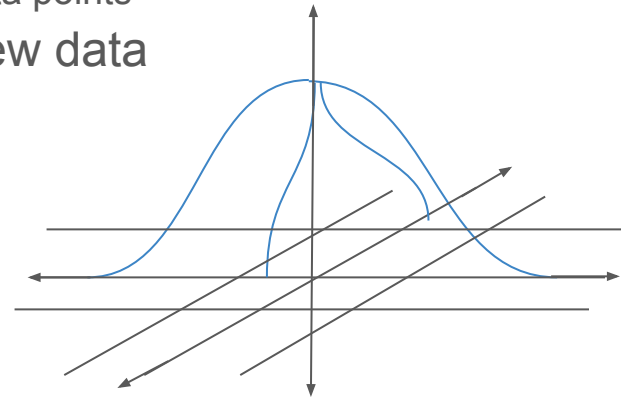
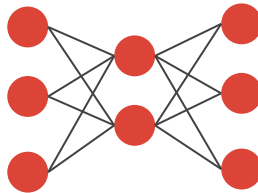


# What is Deep Learning



# Fancy name aside

- Idealized model of biological neurons
- Learn patterns **directly from data**
- Calculus, Linear Algebra, Statistics, Probability, Optimization, Learning Theory, Computational Neuroscience, Systems Neuroscience, Computer Science,...
- Tasks:
  - **Classification** e.g. Given picture a picture, classify what object is in it
  - **Regression** e.g. Given a description of a house, predict the price
  - **Generation** e.g. Learn a distribution and generate data points
- Goal is to learn a function that **generalizes** to new data



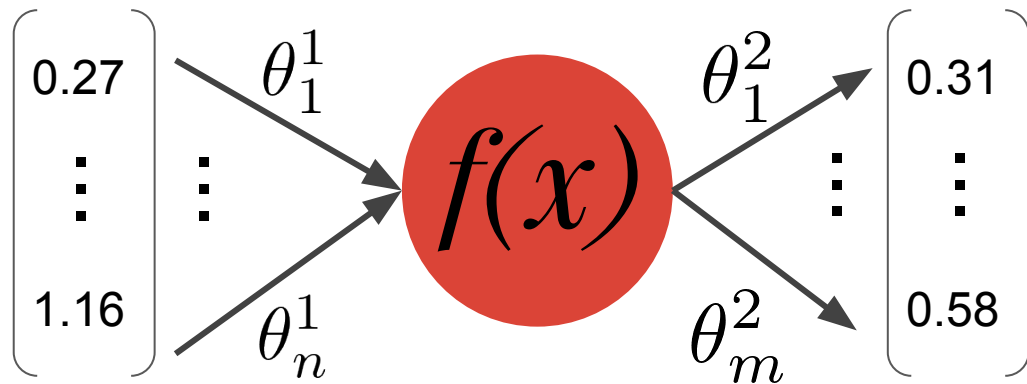
# Artificial Neuron

---

- Nonlinear data
- Classification

# Artificial Neuron

- **Linear Classifier**
- **Linear combination** of input values
- **Nonlinear activation function** of total input
- Each connection has an associated **weight**  $\theta_i^l$





# Nonlinearity

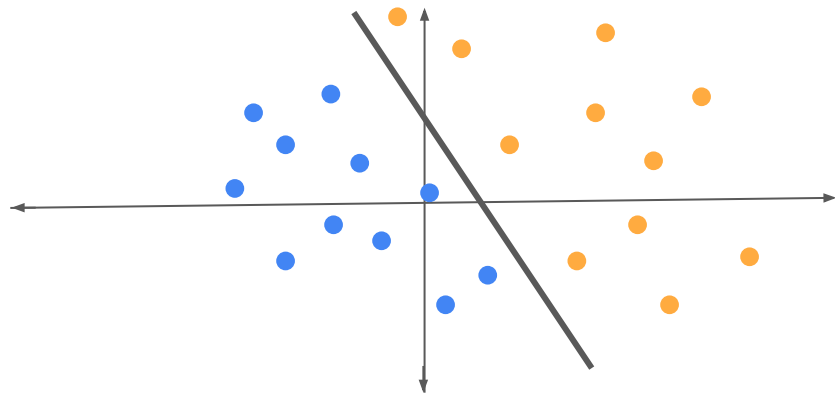
- Pass sum of all inputs ( $x$  in equations below) into a activation function
- Draw nonlinear **decision boundary**
- Handle **nonlinear data**

$$y = Mx + b \rightarrow h(x) = \theta x$$

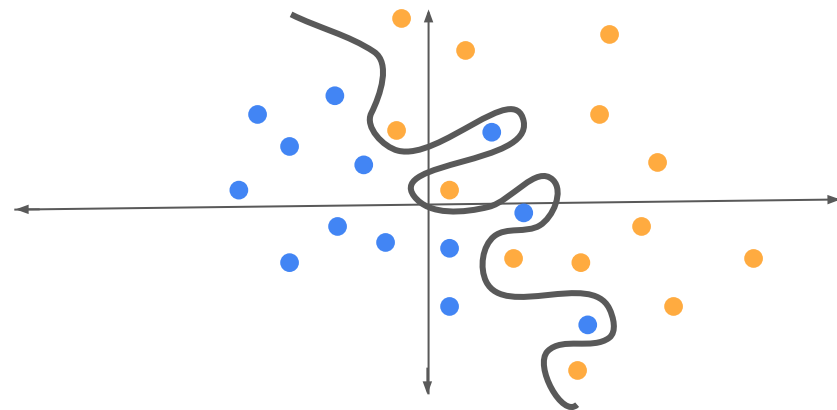


$$f(h(x)) = f(\theta x)$$

Linear Boundary



Nonlinear Boundary

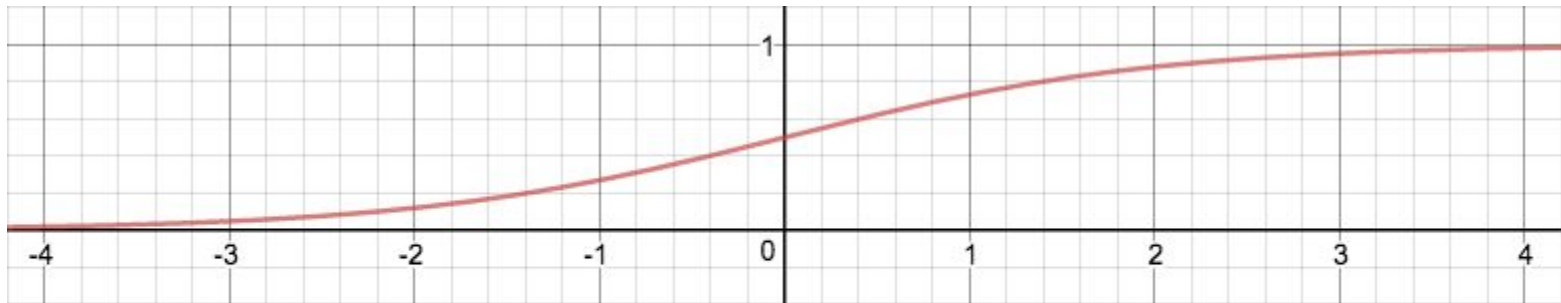


# Logistic Sigmoid

- Used for binary classification
- Outputs value in  $(-1, 1)$
- $f(0) = 0.5$
- Class 1 if output  $< 0.5$
- Class 2 if output  $\geq 0.5$

$$f(x) = \frac{1}{1 + e^{-x}}$$

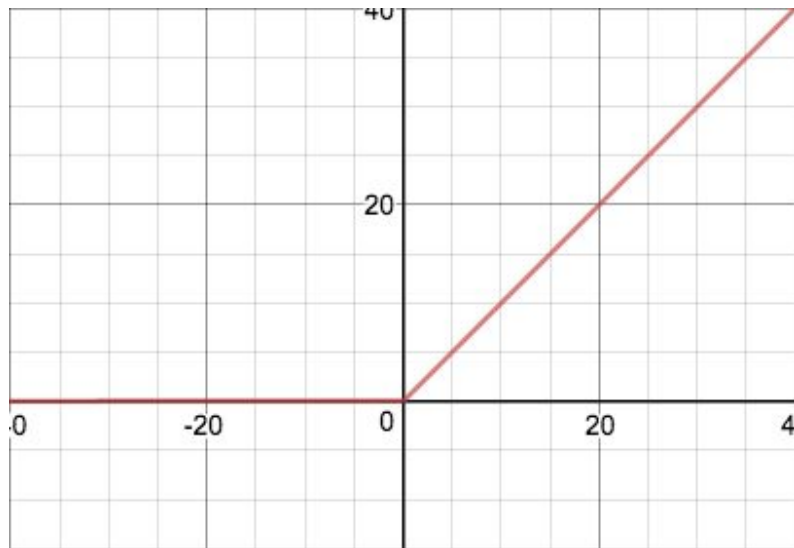
$$f'(x) = \left( \frac{1}{1 + e^{-x}} \right) \left( \frac{e^{-x}}{1 + e^{-x}} \right)$$



# Rectified Linear Unit

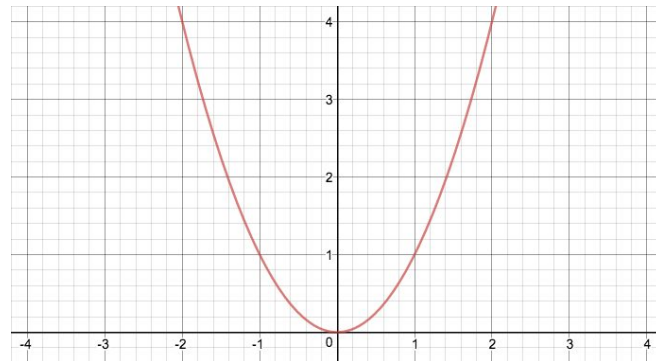
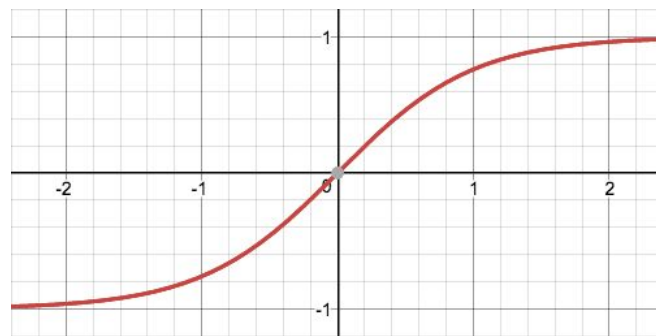
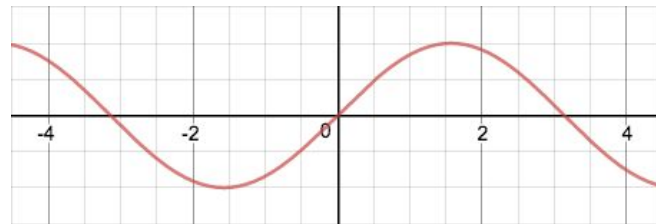
- Currently most popular activation function
- Derivative of 0 is 0 and derivative of x is a constant
- **Inexpensive** to compute
- **Faster** to train <sup>[1]</sup>

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$



# Expressivity

- Expressivity - ability/power to **express solution**
- Possible to use any activation function as long as it's nonlinear, but depends on how you train the network
- Typically will only use **differentiable** functions
- More nonlinearity = more **expressive power**
- Think of activation function as **behavior of a neuron**
  - Biological brain contains various types of neurons with different behaviors

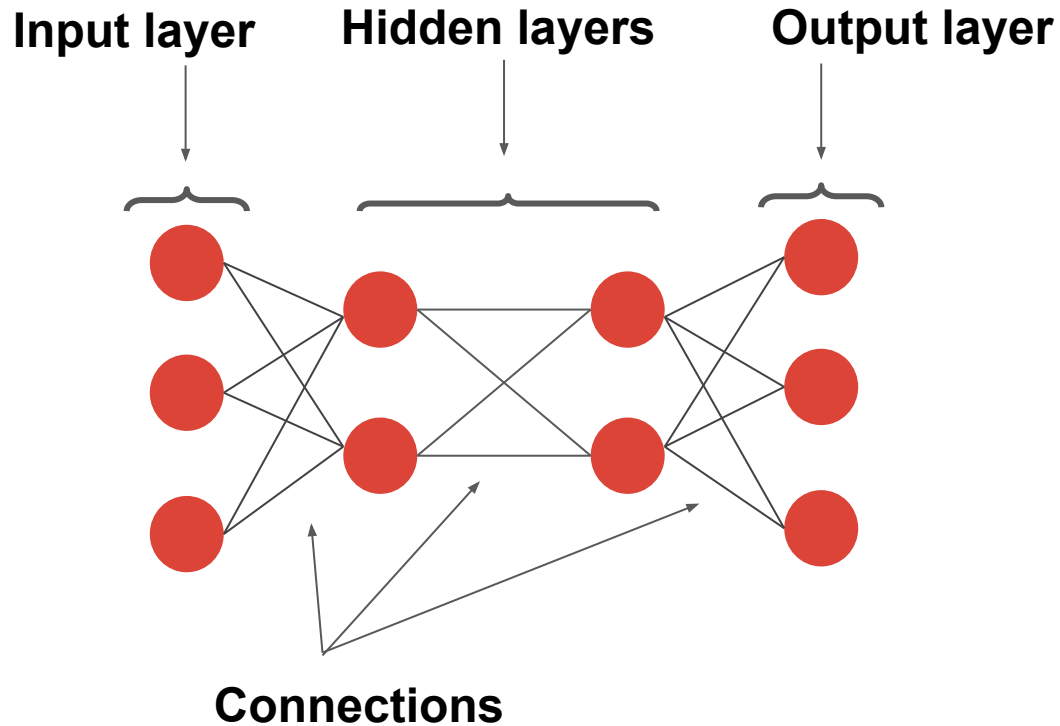


# Artificial Neural Network

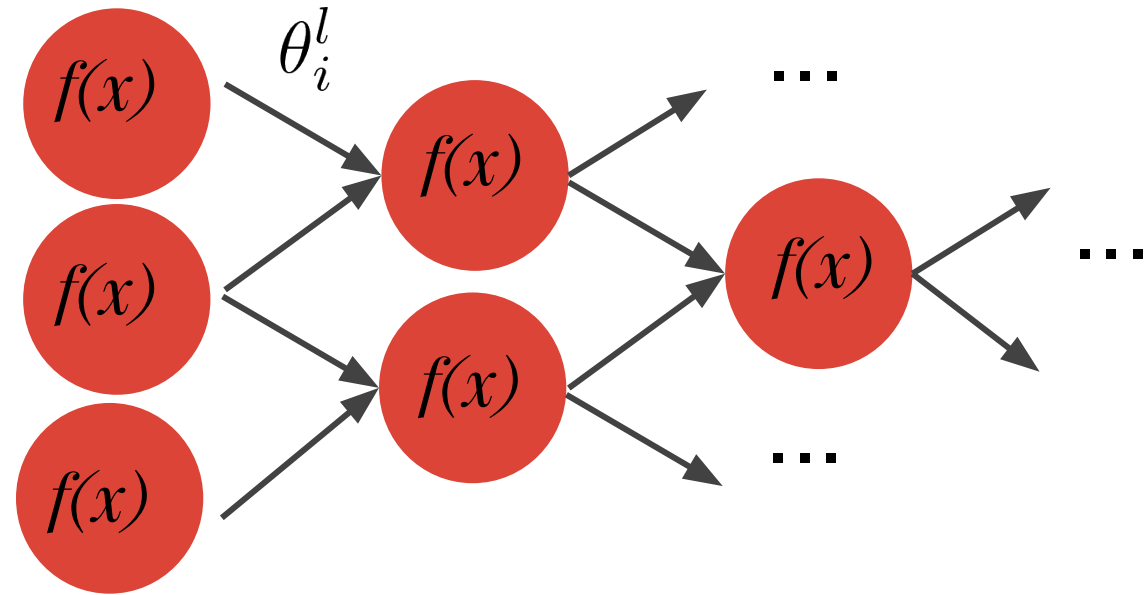
---

- Composition
- Computation Graph
- Deep Neural Networks

# Artificial Neural Network



# Computation Graph



- **Directed computation graph**
- Typically use the same activation function for each node
- Each connection has an associated weight
- **Weight**  $\theta_i^l$  represents how much the receiving neuron should consider it during its own computation

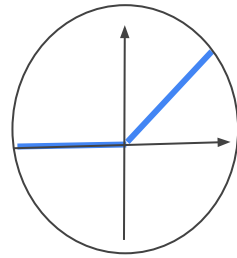
# Ugly Truth

N-dimensional feature vector (Input)

$$\begin{bmatrix} a_{00} & \dots & a_{0n} \\ \vdots & \ddots & \vdots \\ a_{m0} & \dots & a_{mn} \end{bmatrix} \begin{bmatrix} x_0 \\ \vdots \\ x_n \end{bmatrix}$$

Weights between input layer and first hidden layer

Possibly more hidden layers



ReLU

Output layer weights

$$\begin{bmatrix} z_{00} & \dots & z_{0q} \\ \vdots & & \vdots \\ z_{p0} & \dots & z_{pq} \end{bmatrix} \begin{bmatrix} h_0 \\ \vdots \\ h_q \end{bmatrix}_d \longrightarrow \begin{bmatrix} y_0 \\ \vdots \\ y_p \end{bmatrix}$$

Potentially squash last hidden activations through softmax for classification

Output vector



# Universal Approximation Theorem

- Any feed-forward neural network with **at least one hidden layer** using **logistic sigmoid** activation function can **approximate** any **continuous function** over a **compact set** to an arbitrary accuracy given sufficient parameters

(Cybenko, 1989) <sup>[2]</sup>

- Any feed-forward neural network with **at least one hidden layer** using **any activation function** can **approximate** any **continuous function** over a **compact set** to an arbitrary accuracy given sufficient parameters

(Hornik, 1991) <sup>[3]</sup>

# How do they actually learn?

---

- Costs
- Gradients
- Error corrections

# Dataset

Split - percentages depend on available data

TRAINING  
~70%

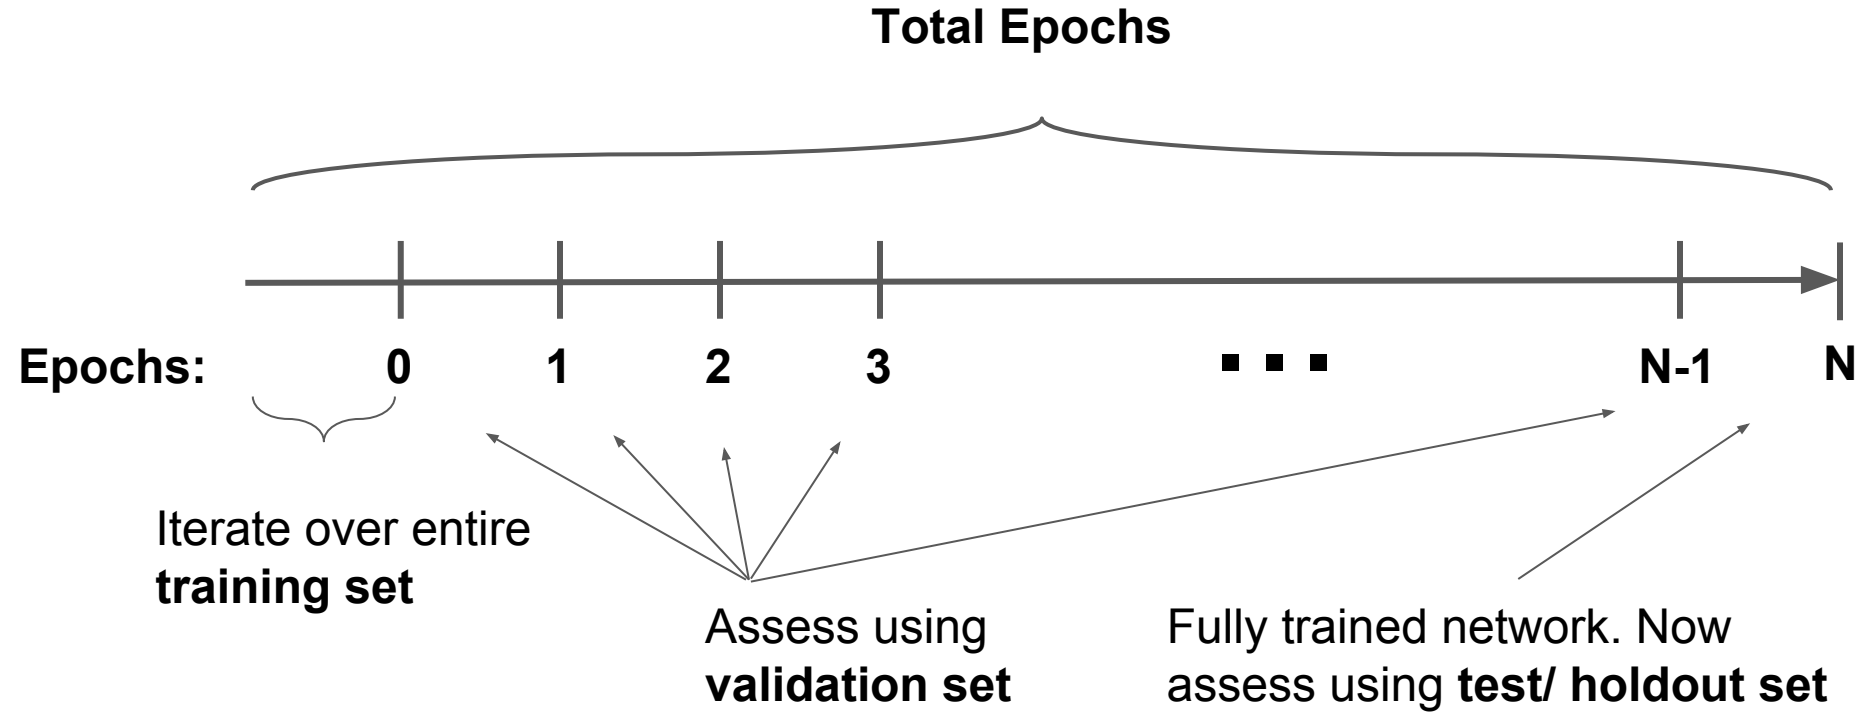
VALIDATION  
~10%

TEST/  
HOLDOUT  
~20%

Balance - data distributed close to uniformly



# Training



# Cost Function

- Also known as **objective** function, **loss** function
- Tells learning algorithm how well/poorly it's doing
- Task-dependent hyperparameter

## Least Square Error

for regression

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

## Cross-Entropy

for classification

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})))$$

# Gradient Descent

- Minimize cost function w.r.t. Weights
- Each update is also known as a **Gradient Step**

Also written  $\nabla J(\theta)$

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Scalar learning rate

Individual weights

Cost/objective/loss function

Vector of weights

The diagram illustrates the gradient descent update equation:  $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$ . Annotations include: 'Scalar learning rate' pointing to  $\alpha$ ; 'Individual weights' pointing to the  $\theta_j$  on the left; 'Cost/objective/loss function' pointing to  $J(\theta)$ ; 'Vector of weights' pointing to  $\theta$  inside  $J(\theta)$ ; and 'Also written  $\nabla J(\theta)$ ' with a bracket pointing to the derivative term  $\frac{\partial}{\partial \theta_j} J(\theta)$ .

# Gradient Descent cont.

## Batch/Vanilla Gradient Descent

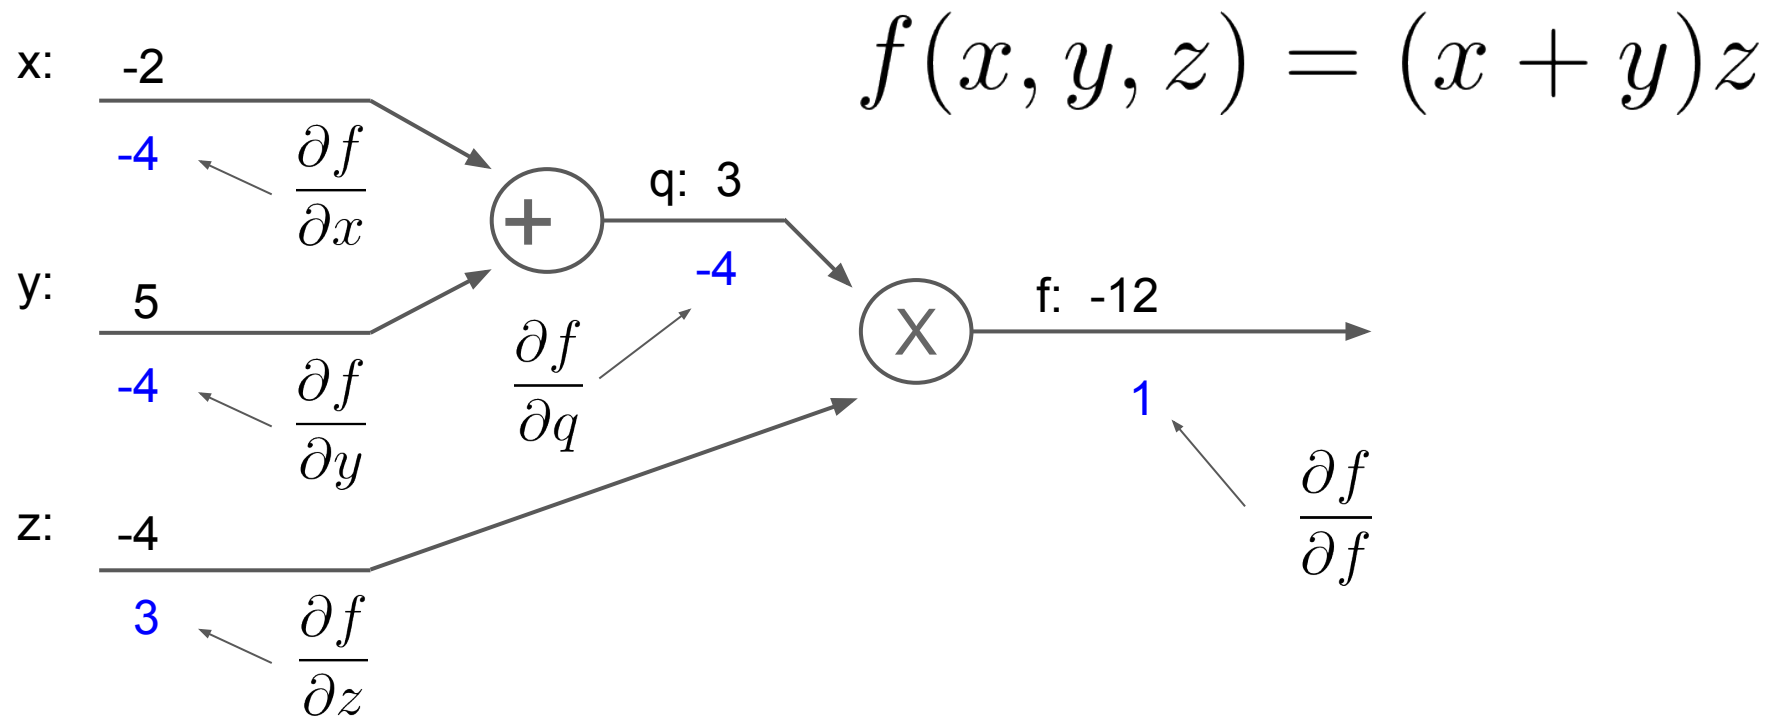
- Feed in entire dataset and then make one update to weights
- Slow and must store dataset in memory

## Stochastic Gradient Descent

- Feed in one example and then update weights
- Loop over dataset
- Fast and implicitly regularizes

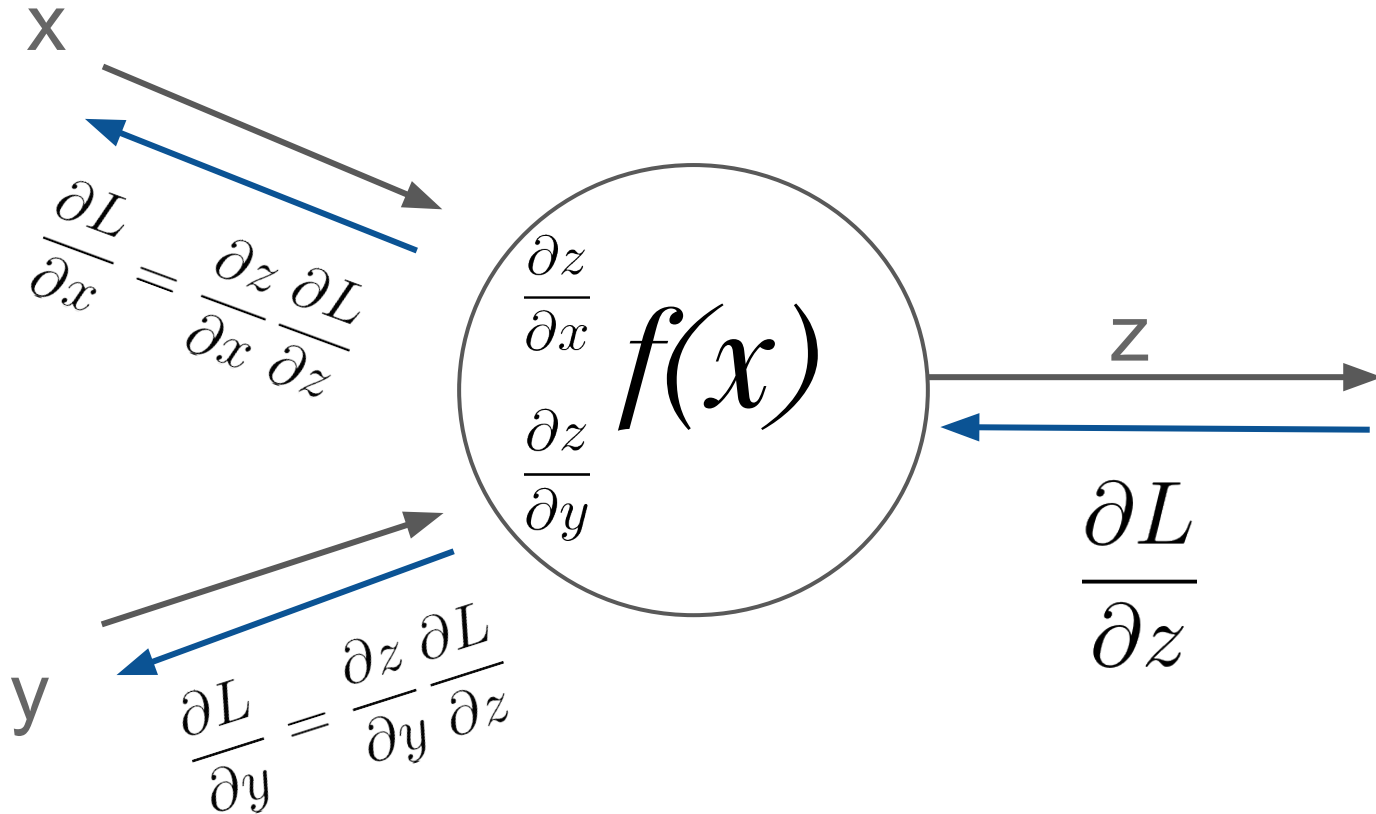
# Backpropagation

- Compute gradient of loss w.r.t. to weights





# Backpropagation



# Reminders

## 1. **Stochastic Gradient Methods for Large Scale Machine Learning**

Location: Harvard, Maxwell Dworkin G115

Time: 4 PM, ice cream social at 3:30

## 2. **MIT MIC tomorrow at 5 PM on CycleGAN**

Location: 56-154

Time: 5 PM

Food will be provided

# References & Further Reading

- [0] Le, Quoc V. "A Tutorial on Deep Learning Part 1: Nonlinear Classifiers and The Backpropagation Algorithm." (2015).
- [1] LeCun, Yann, Yoshua Bengio, and Geoffrey Hinton. "Deep learning." *Nature* 521.7553 (2015): 436-444.
- [2] Gybenko, G. "Approximation by superposition of sigmoidal functions." *Mathematics of Control, Signals and Systems* 2.4 (1989): 303-314.
- [3] Hornik, Kurt. "Approximation capabilities of multilayer feedforward networks." *Neural networks* 4.2 (1991): 251-257.
- [4] Stanford CS231n: Convolutional Neural Networks for Visual Recognition
- [5] <https://github.com/ch3njust1n/bestofml>